

int skills Chris C

August 27, 2018

1 int skills makeup work

Christopher Csiszar

1.0.1 hyperparameter tuning

I haven't worked with linear regressors in a while, so I'll be practicing on those + a simple RF.

Note the linear regressors below have a built in cross validation function, however I will be returning to basics and tuning the models in a simple way.

```
In [77]: import pandas as pd
import numpy as np
import datetime
from datetime import datetime
import dateutil.parser
from sklearn.model_selection import train_test_split

from sklearn import datasets, linear_model
from sklearn.metrics import mean_squared_error, r2_score

from sklearn.linear_model import Ridge
from sklearn.ensemble import RandomForestRegressor
```

```
In [35]: ! pwd
```

```
/Users/chrispaul/Desktop/classes/intskills/kechup
```

```
In [36]: train = pd.read_csv('Train.csv')
```

```
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/IPython/core/interactiveshell.py
interactivity=interactivity, compiler=compiler, result=result)
```

```
In [37]: train = train[0:23900]
```

```
In [38]: train.tail()
```

```

Out [38]:
      SalesID  SalePrice  under_20k  MachineID  ModelID  datasource  \
23895  1222558.0    81000.0         0.0  1019939.0   14287.0        121.0
23896  1222569.0    23000.0         0.0  1057622.0    3350.0        121.0
23897  1222570.0    17000.0         1.0  1001012.0    3414.0        121.0
23898  1222571.0    16000.0         1.0  1023061.0    1528.0        121.0
23899  1222575.0    48000.0         0.0  1061091.0   28920.0        121.0

      auctioneerID  YearMade  MachineHoursCurrentMeter  saledate  \
23895           3.0    2006.0                   3395.0  12/22/11 0:00
23896           3.0    1000.0                   0.0  12/22/11 0:00
23897           3.0    1000.0                  13507.0  12/22/11 0:00
23898           3.0    1000.0                   5417.0  12/22/11 0:00
23899           3.0    2005.0                   0.0  12/22/11 0:00

      state ProductGroup  Enclosure
23895  Virginia         TEX  EROPS w AC
23896  Virginia         MG   EROPS
23897  Virginia         TEX  EROPS
23898  Virginia         TTT   OROPS
23899  Virginia         WL  EROPS w AC

```

```

In [39]: train.describe(include='all')

```

```

Out [39]:
      SalesID  SalePrice  under_20k  MachineID  ModelID  \
count  2.390000e+04  23900.000000  23900.000000  2.390000e+04  23900.000000
unique          NaN          NaN          NaN          NaN          NaN
top           NaN          NaN          NaN          NaN          NaN
freq          NaN          NaN          NaN          NaN          NaN
mean  1.179777e+06  33888.044477    0.375941  8.563724e+05  8326.160795
std  2.392394e+04  24591.558640    0.484375  3.120474e+05  7562.924092
min  1.139246e+06   4750.000000    0.000000  3.230000e+02   28.000000
25%  1.159416e+06  15000.000000    0.000000  7.718310e+05  3357.000000
50%  1.178196e+06  26500.000000    0.000000  1.018198e+06  4666.000000
75%  1.200834e+06  45625.000000    1.000000  1.043985e+06  13395.000000
max  1.222575e+06  142000.000000    1.000000  1.069977e+06  37198.000000

      datasource  auctioneerID  YearMade  MachineHoursCurrentMeter  \
count  23900.0    23900.0  23900.000000    23900.000000
unique    NaN      NaN      NaN      NaN
top      NaN      NaN      NaN      NaN
freq      NaN      NaN      NaN      NaN
mean    121.0      3.0  1821.906444    4649.534142
std      0.0      0.0   384.023768    6066.597357
min    121.0      3.0  1000.000000      0.000000
25%    121.0      3.0  1996.000000    930.000000
50%    121.0      3.0  2001.000000   2638.000000
75%    121.0      3.0  2004.000000   6470.000000
max    121.0      3.0  2010.000000  220893.000000

```

	saledate	state	ProductGroup	Enclosure
count	23900	23900	23900	23884
unique	823	52	6	3
top	1/29/09 0:00	Texas	TEX	OROPS
freq	200	2823	6579	9399
mean	NaN	NaN	NaN	NaN
std	NaN	NaN	NaN	NaN
min	NaN	NaN	NaN	NaN
25%	NaN	NaN	NaN	NaN
50%	NaN	NaN	NaN	NaN
75%	NaN	NaN	NaN	NaN
max	NaN	NaN	NaN	NaN

In [40]: train['datesale'] = train.apply(lambda row: dateutil.parser.parse(row['saledate']).year,

In [41]: train

Out[41]:

	SalesID	SalePrice	under_20k	MachineID	ModelID	datasource	\
0	1139246.0	66000.0	0.0	999089.0	3157.0	121.0	
1	1139248.0	57000.0	0.0	117657.0	77.0	121.0	
2	1139249.0	10000.0	1.0	434808.0	7009.0	121.0	
3	1139251.0	38500.0	0.0	1026470.0	332.0	121.0	
4	1139253.0	11000.0	1.0	1057373.0	17311.0	121.0	
5	1139255.0	26500.0	0.0	1001274.0	4605.0	121.0	
6	1139256.0	21000.0	0.0	772701.0	1937.0	121.0	
7	1139261.0	27000.0	0.0	902002.0	3539.0	121.0	
8	1139272.0	21500.0	0.0	1036251.0	36003.0	121.0	
9	1139275.0	65000.0	0.0	1016474.0	3883.0	121.0	
10	1139278.0	24000.0	0.0	1024998.0	4605.0	121.0	
11	1139282.0	22500.0	0.0	319906.0	5255.0	121.0	
12	1139283.0	36000.0	0.0	1052214.0	2232.0	121.0	
13	1139284.0	30500.0	0.0	1068082.0	3542.0	121.0	
14	1139290.0	28000.0	0.0	1058450.0	5162.0	121.0	
15	1139291.0	19000.0	1.0	1004810.0	4604.0	121.0	
16	1139292.0	13500.0	1.0	1026973.0	9510.0	121.0	
17	1139299.0	9500.0	1.0	1002713.0	21442.0	121.0	
18	1139301.0	12500.0	1.0	125790.0	7040.0	121.0	
19	1139304.0	11500.0	1.0	1011914.0	3177.0	121.0	
20	1139311.0	41000.0	0.0	1014135.0	8867.0	121.0	
21	1139333.0	34500.0	0.0	999192.0	3350.0	121.0	
22	1139344.0	26000.0	0.0	1044500.0	7040.0	121.0	
23	1139346.0	73000.0	0.0	821452.0	85.0	121.0	
24	1139348.0	33000.0	0.0	294562.0	3542.0	121.0	
25	1139351.0	12500.0	1.0	833838.0	7009.0	121.0	
26	1139354.0	15500.0	1.0	565440.0	7040.0	121.0	
27	1139356.0	53000.0	0.0	1004127.0	25458.0	121.0	
28	1139357.0	46000.0	0.0	44800.0	19167.0	121.0	

29	1139358.0	89000.0	0.0	1018076.0	1333.0	121.0
...
23870	1222464.0	27000.0	0.0	1027624.0	328.0	121.0
23871	1222466.0	79000.0	0.0	1008813.0	16506.0	121.0
23872	1222468.0	85000.0	0.0	520588.0	23926.0	121.0
23873	1222471.0	27000.0	0.0	1050702.0	22155.0	121.0
23874	1222474.0	7000.0	1.0	1022467.0	18263.0	121.0
23875	1222505.0	70000.0	0.0	198296.0	1263.0	121.0
23876	1222507.0	36000.0	0.0	213148.0	3542.0	121.0
23877	1222509.0	12500.0	1.0	1063187.0	4107.0	121.0
23878	1222510.0	26000.0	0.0	1065790.0	23737.0	121.0
23879	1222511.0	20000.0	0.0	1046116.0	28587.0	121.0
23880	1222512.0	11500.0	1.0	1033783.0	5436.0	121.0
23881	1222514.0	8500.0	1.0	1032478.0	3170.0	121.0
23882	1222516.0	10000.0	1.0	1057872.0	3170.0	121.0
23883	1222531.0	5500.0	1.0	1039959.0	17592.0	121.0
23884	1222534.0	12000.0	1.0	1008217.0	1958.0	121.0
23885	1222537.0	18000.0	1.0	1025370.0	3883.0	121.0
23886	1222538.0	23500.0	0.0	1046410.0	3883.0	121.0
23887	1222540.0	24000.0	0.0	1054036.0	3893.0	121.0
23888	1222541.0	13000.0	1.0	1055823.0	3369.0	121.0
23889	1222542.0	63000.0	0.0	753302.0	3886.0	121.0
23890	1222543.0	50000.0	0.0	1049123.0	3886.0	121.0
23891	1222544.0	12000.0	1.0	1002257.0	3883.0	121.0
23892	1222551.0	29500.0	0.0	705473.0	3539.0	121.0
23893	1222552.0	21000.0	0.0	1040718.0	18110.0	121.0
23894	1222553.0	29250.0	0.0	1022899.0	22854.0	121.0
23895	1222558.0	81000.0	0.0	1019939.0	14287.0	121.0
23896	1222569.0	23000.0	0.0	1057622.0	3350.0	121.0
23897	1222570.0	17000.0	1.0	1001012.0	3414.0	121.0
23898	1222571.0	16000.0	1.0	1023061.0	1528.0	121.0
23899	1222575.0	48000.0	0.0	1061091.0	28920.0	121.0

	auctioneerID	YearMade	MachineHoursCurrentMeter	saledate \
0	3.0	2004.0	68.0	11/16/06 0:00
1	3.0	1996.0	4640.0	3/26/04 0:00
2	3.0	2001.0	2838.0	2/26/04 0:00
3	3.0	2001.0	3486.0	5/19/11 0:00
4	3.0	2007.0	722.0	7/23/09 0:00
5	3.0	2004.0	508.0	12/18/08 0:00
6	3.0	1993.0	11540.0	8/26/04 0:00
7	3.0	2001.0	4883.0	11/17/05 0:00
8	3.0	2008.0	302.0	8/27/09 0:00
9	3.0	1000.0	20700.0	8/9/07 0:00
10	3.0	2004.0	1414.0	8/21/08 0:00
11	3.0	1998.0	2764.0	8/24/06 0:00
12	3.0	1998.0	0.0	10/20/05 0:00
13	3.0	2001.0	1921.0	1/26/06 0:00

14	3.0	2004.0	320.0	1/3/06	0:00
15	3.0	1999.0	2450.0	11/16/06	0:00
16	3.0	1999.0	1972.0	6/14/07	0:00
17	3.0	2003.0	0.0	1/28/10	0:00
18	3.0	2001.0	994.0	3/9/06	0:00
19	3.0	1991.0	8005.0	11/17/05	0:00
20	3.0	2000.0	3259.0	5/18/06	0:00
21	3.0	1000.0	16328.0	10/19/06	0:00
22	3.0	2005.0	109.0	10/25/07	0:00
23	3.0	1996.0	17033.0	10/19/06	0:00
24	3.0	2001.0	1877.0	5/20/04	0:00
25	3.0	2003.0	1028.0	3/9/06	0:00
26	3.0	2003.0	356.0	3/9/06	0:00
27	3.0	2000.0	0.0	2/22/07	0:00
28	3.0	2004.0	904.0	8/9/07	0:00
29	3.0	1998.0	10466.0	6/1/06	0:00
...
23870	3.0	1000.0	9139.0	12/14/11	0:00
23871	3.0	2008.0	5234.0	12/15/11	0:00
23872	3.0	2000.0	11657.0	12/22/11	0:00
23873	3.0	2005.0	0.0	12/22/11	0:00
23874	3.0	2005.0	1032.0	12/22/11	0:00
23875	3.0	2001.0	11589.0	12/22/11	0:00
23876	3.0	2005.0	8563.0	12/22/11	0:00
23877	3.0	1000.0	9229.0	12/22/11	0:00
23878	3.0	1000.0	15060.0	12/22/11	0:00
23879	3.0	1000.0	17674.0	12/22/11	0:00
23880	3.0	1997.0	1533.0	12/22/11	0:00
23881	3.0	1000.0	881.0	12/22/11	0:00
23882	3.0	1000.0	2030.0	12/22/11	0:00
23883	3.0	1999.0	3333.0	12/22/11	0:00
23884	3.0	1000.0	5994.0	12/22/11	0:00
23885	3.0	1000.0	29598.0	12/22/11	0:00
23886	3.0	1000.0	10178.0	12/22/11	0:00
23887	3.0	1000.0	4607.0	12/22/11	0:00
23888	3.0	1000.0	0.0	12/22/11	0:00
23889	3.0	1997.0	23386.0	12/22/11	0:00
23890	3.0	1000.0	34057.0	12/22/11	0:00
23891	3.0	1000.0	24074.0	12/22/11	0:00
23892	3.0	2002.0	3138.0	12/21/11	0:00
23893	3.0	2007.0	1290.0	12/16/11	0:00
23894	3.0	2004.0	2299.0	12/16/11	0:00
23895	3.0	2006.0	3395.0	12/22/11	0:00
23896	3.0	1000.0	0.0	12/22/11	0:00
23897	3.0	1000.0	13507.0	12/22/11	0:00
23898	3.0	1000.0	5417.0	12/22/11	0:00
23899	3.0	2005.0	0.0	12/22/11	0:00

	state	ProductGroup	Enclosure	datesale
0	Alabama	WL	EROPS w AC	2006
1	North Carolina	WL	EROPS w AC	2004
2	New York	SSL	OROPS	2004
3	Texas	TEX	EROPS w AC	2011
4	New York	SSL	EROPS	2009
5	Arizona	BL	OROPS	2008
6	Florida	TEX	EROPS	2004
7	Illinois	BL	OROPS	2005
8	Texas	TEX	EROPS	2009
9	Florida	WL	EROPS w AC	2007
10	Oregon	BL	OROPS	2008
11	Ohio	TTT	EROPS w AC	2006
12	Ohio	TEX	EROPS	2005
13	Texas	BL	OROPS	2006
14	North Carolina	BL	OROPS	2006
15	Arkansas	BL	OROPS	2006
16	Florida	TEX	EROPS	2007
17	Wisconsin	TEX	EROPS	2010
18	North Carolina	TEX	EROPS	2006
19	Illinois	BL	EROPS	2005
20	Kansas	TEX	EROPS w AC	2006
21	Nevada	MG	EROPS	2006
22	Iowa	TEX	EROPS w AC	2007
23	Maine	WL	EROPS w AC	2006
24	Texas	BL	OROPS	2004
25	Massachusetts	SSL	EROPS	2006
26	California	TEX	EROPS	2006
27	Texas	TEX	EROPS w AC	2007
28	Louisiana	MG	OROPS	2007
29	Minnesota	TEX	EROPS w AC	2006
...
23870	Washington	TEX	EROPS	2011
23871	Arkansas	TEX	EROPS w AC	2011
23872	Florida	MG	EROPS w AC	2011
23873	Florida	TTT	OROPS	2011
23874	Maryland	SSL	OROPS	2011
23875	New Mexico	TEX	EROPS w AC	2011
23876	New Mexico	BL	EROPS w AC	2011
23877	Ohio	TTT	OROPS	2011
23878	Ohio	TEX	EROPS	2011
23879	Ohio	TEX	EROPS	2011
23880	Ohio	TEX	EROPS	2011
23881	Ohio	BL	EROPS	2011
23882	Ohio	BL	EROPS	2011
23883	North Carolina	SSL	OROPS	2011
23884	California	TEX	EROPS	2011
23885	California	WL	EROPS	2011

23886	California	WL	EROPS	2011
23887	California	WL	EROPS	2011
23888	California	MG	EROPS	2011
23889	California	WL	EROPS w AC	2011
23890	California	WL	EROPS w AC	2011
23891	California	WL	EROPS w AC	2011
23892	California	BL	EROPS w AC	2011
23893	Florida	BL	OROPS	2011
23894	Florida	WL	OROPS	2011
23895	Virginia	TEX	EROPS w AC	2011
23896	Virginia	MG	EROPS	2011
23897	Virginia	TEX	EROPS	2011
23898	Virginia	TTT	OROPS	2011
23899	Virginia	WL	EROPS w AC	2011

[23900 rows x 14 columns]

```
In [42]: train['yearsold'] = train.datesale - train.YearMade
```

```
In [43]: train.head()
```

```
Out [43]:
```

	SalesID	SalePrice	under_20k	MachineID	ModelID	datasource	\
0	1139246.0	66000.0	0.0	999089.0	3157.0	121.0	
1	1139248.0	57000.0	0.0	117657.0	77.0	121.0	
2	1139249.0	10000.0	1.0	434808.0	7009.0	121.0	
3	1139251.0	38500.0	0.0	1026470.0	332.0	121.0	
4	1139253.0	11000.0	1.0	1057373.0	17311.0	121.0	

	auctioneerID	YearMade	MachineHoursCurrentMeter	saledate	\
0	3.0	2004.0		68.0 11/16/06 0:00	
1	3.0	1996.0		4640.0 3/26/04 0:00	
2	3.0	2001.0		2838.0 2/26/04 0:00	
3	3.0	2001.0		3486.0 5/19/11 0:00	
4	3.0	2007.0		722.0 7/23/09 0:00	

	state	ProductGroup	Enclosure	datesale	yearsold
0	Alabama	WL	EROPS w AC	2006	2.0
1	North Carolina	WL	EROPS w AC	2004	8.0
2	New York	SSL	OROPS	2004	3.0
3	Texas	TEX	EROPS w AC	2011	10.0
4	New York	SSL	EROPS	2009	2.0

```
In [44]: train2 = train.drop(['saledate', "auctioneerID"], axis=1)
```

```
In [45]: train2.head()
```

```
Out [45]:
```

	SalesID	SalePrice	under_20k	MachineID	ModelID	datasource	YearMade	\
0	1139246.0	66000.0	0.0	999089.0	3157.0	121.0	2004.0	
1	1139248.0	57000.0	0.0	117657.0	77.0	121.0	1996.0	

2	1139249.0	10000.0	1.0	434808.0	7009.0	121.0	2001.0
3	1139251.0	38500.0	0.0	1026470.0	332.0	121.0	2001.0
4	1139253.0	11000.0	1.0	1057373.0	17311.0	121.0	2007.0

	MachineHoursCurrentMeter	state	ProductGroup	Enclosure	\
0	68.0	Alabama	WL	EROPS w AC	
1	4640.0	North Carolina	WL	EROPS w AC	
2	2838.0	New York	SSL	OROPS	
3	3486.0	Texas	TEX	EROPS w AC	
4	722.0	New York	SSL	EROPS	

	datesale	yearsold
0	2006	2.0
1	2004	8.0
2	2004	3.0
3	2011	10.0
4	2009	2.0

```
In [46]: def proc_col(col, train_col=None):
        """Encodes a pandas column with continous ids.
        """
        if train_col is not None:
            uniq = train_col.unique()
        else:
            uniq = col.unique()
        name2idx = {o:i for i,o in enumerate(uniq)}
        return name2idx, np.array([name2idx.get(x, -1) for x in col]), len(uniq)
```

```
In [47]: def encode_data(df, train=None):
        """ Encodes rating data with continous user and movie ids.
        If train is provided, encodes df with the same encoding as train.
        """
        df = df.copy()
        for col_name in ["state", "ProductGroup", "Enclosure"]:
            train_col = None
            if train is not None:
                train_col = train[col_name]
            _,col,_ = proc_col(df[col_name], train_col)
            df[col_name] = col
            df = df[df[col_name] >= 0]
        return df
```

```
In [48]: data = encode_data(train2)
```

```
In [49]: data.head()
```

```
Out[49]:
```

	SalesID	SalePrice	under_20k	MachineID	ModelID	datasource	YearMade	\
0	1139246.0	66000.0	0.0	999089.0	3157.0	121.0	2004.0	
1	1139248.0	57000.0	0.0	117657.0	77.0	121.0	1996.0	

2	1139249.0	10000.0	1.0	434808.0	7009.0	121.0	2001.0
3	1139251.0	38500.0	0.0	1026470.0	332.0	121.0	2001.0
4	1139253.0	11000.0	1.0	1057373.0	17311.0	121.0	2007.0

	MachineHoursCurrentMeter	state	ProductGroup	Enclosure	datesale	\
0	68.0	0	0	0	2006	
1	4640.0	1	0	0	2004	
2	2838.0	2	1	1	2004	
3	3486.0	3	2	0	2011	
4	722.0	2	1	2	2009	

	yearsold
0	2.0
1	8.0
2	3.0
3	10.0
4	2.0

In [50]: data.describe(include='all')

Out [50]:

	SalesID	SalePrice	under_20k	MachineID	ModelID	\
count	2.390000e+04	23900.000000	23900.000000	2.390000e+04	23900.000000	
mean	1.179777e+06	33888.044477	0.375941	8.563724e+05	8326.160795	
std	2.392394e+04	24591.558640	0.484375	3.120474e+05	7562.924092	
min	1.139246e+06	4750.000000	0.000000	3.230000e+02	28.000000	
25%	1.159416e+06	15000.000000	0.000000	7.718310e+05	3357.000000	
50%	1.178196e+06	26500.000000	0.000000	1.018198e+06	4666.000000	
75%	1.200834e+06	45625.000000	1.000000	1.043985e+06	13395.000000	
max	1.222575e+06	142000.000000	1.000000	1.069977e+06	37198.000000	

	datasource	YearMade	MachineHoursCurrentMeter	state	\
count	23900.0	23900.000000	23900.000000	23900.000000	
mean	121.0	1821.906444	4649.534142	16.290084	
std	0.0	384.023768	6066.597357	13.107064	
min	121.0	1000.000000	0.000000	0.000000	
25%	121.0	1996.000000	930.000000	5.000000	
50%	121.0	2001.000000	2638.000000	16.000000	
75%	121.0	2004.000000	6470.000000	26.000000	
max	121.0	2010.000000	220893.000000	51.000000	

	ProductGroup	Enclosure	datesale	yearsold
count	23900.000000	23900.000000	23900.000000	23900.000000
mean	2.184100	0.858536	2008.424477	186.518033
std	1.487749	0.767305	2.121588	384.468595
min	0.000000	0.000000	2004.000000	0.000000
25%	1.000000	0.000000	2007.000000	4.000000
50%	2.000000	1.000000	2009.000000	7.000000
75%	3.000000	1.000000	2010.000000	13.000000
max	5.000000	3.000000	2011.000000	1011.000000

2 Regression

```
In [56]: y = data.SalePrice
```

```
In [57]: x = data.drop(['SalePrice', "under_20k"], axis=1)
```

```
In [58]: X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=
```

```
In [59]: len(X_train)
```

```
Out [59]: 19120
```

```
In [60]: x
```

```
Out [60]:
```

	SalesID	MachineID	ModelID	datasource	YearMade	\
0	1139246.0	999089.0	3157.0	121.0	2004.0	
1	1139248.0	117657.0	77.0	121.0	1996.0	
2	1139249.0	434808.0	7009.0	121.0	2001.0	
3	1139251.0	1026470.0	332.0	121.0	2001.0	
4	1139253.0	1057373.0	17311.0	121.0	2007.0	
5	1139255.0	1001274.0	4605.0	121.0	2004.0	
6	1139256.0	772701.0	1937.0	121.0	1993.0	
7	1139261.0	902002.0	3539.0	121.0	2001.0	
8	1139272.0	1036251.0	36003.0	121.0	2008.0	
9	1139275.0	1016474.0	3883.0	121.0	1000.0	
10	1139278.0	1024998.0	4605.0	121.0	2004.0	
11	1139282.0	319906.0	5255.0	121.0	1998.0	
12	1139283.0	1052214.0	2232.0	121.0	1998.0	
13	1139284.0	1068082.0	3542.0	121.0	2001.0	
14	1139290.0	1058450.0	5162.0	121.0	2004.0	
15	1139291.0	1004810.0	4604.0	121.0	1999.0	
16	1139292.0	1026973.0	9510.0	121.0	1999.0	
17	1139299.0	1002713.0	21442.0	121.0	2003.0	
18	1139301.0	125790.0	7040.0	121.0	2001.0	
19	1139304.0	1011914.0	3177.0	121.0	1991.0	
20	1139311.0	1014135.0	8867.0	121.0	2000.0	
21	1139333.0	999192.0	3350.0	121.0	1000.0	
22	1139344.0	1044500.0	7040.0	121.0	2005.0	
23	1139346.0	821452.0	85.0	121.0	1996.0	
24	1139348.0	294562.0	3542.0	121.0	2001.0	
25	1139351.0	833838.0	7009.0	121.0	2003.0	
26	1139354.0	565440.0	7040.0	121.0	2003.0	
27	1139356.0	1004127.0	25458.0	121.0	2000.0	
28	1139357.0	44800.0	19167.0	121.0	2004.0	
29	1139358.0	1018076.0	1333.0	121.0	1998.0	
...	
23870	1222464.0	1027624.0	328.0	121.0	1000.0	
23871	1222466.0	1008813.0	16506.0	121.0	2008.0	
23872	1222468.0	520588.0	23926.0	121.0	2000.0	

23873	1222471.0	1050702.0	22155.0	121.0	2005.0
23874	1222474.0	1022467.0	18263.0	121.0	2005.0
23875	1222505.0	198296.0	1263.0	121.0	2001.0
23876	1222507.0	213148.0	3542.0	121.0	2005.0
23877	1222509.0	1063187.0	4107.0	121.0	1000.0
23878	1222510.0	1065790.0	23737.0	121.0	1000.0
23879	1222511.0	1046116.0	28587.0	121.0	1000.0
23880	1222512.0	1033783.0	5436.0	121.0	1997.0
23881	1222514.0	1032478.0	3170.0	121.0	1000.0
23882	1222516.0	1057872.0	3170.0	121.0	1000.0
23883	1222531.0	1039959.0	17592.0	121.0	1999.0
23884	1222534.0	1008217.0	1958.0	121.0	1000.0
23885	1222537.0	1025370.0	3883.0	121.0	1000.0
23886	1222538.0	1046410.0	3883.0	121.0	1000.0
23887	1222540.0	1054036.0	3893.0	121.0	1000.0
23888	1222541.0	1055823.0	3369.0	121.0	1000.0
23889	1222542.0	753302.0	3886.0	121.0	1997.0
23890	1222543.0	1049123.0	3886.0	121.0	1000.0
23891	1222544.0	1002257.0	3883.0	121.0	1000.0
23892	1222551.0	705473.0	3539.0	121.0	2002.0
23893	1222552.0	1040718.0	18110.0	121.0	2007.0
23894	1222553.0	1022899.0	22854.0	121.0	2004.0
23895	1222558.0	1019939.0	14287.0	121.0	2006.0
23896	1222569.0	1057622.0	3350.0	121.0	1000.0
23897	1222570.0	1001012.0	3414.0	121.0	1000.0
23898	1222571.0	1023061.0	1528.0	121.0	1000.0
23899	1222575.0	1061091.0	28920.0	121.0	2005.0

	MachineHoursCurrentMeter	state	ProductGroup	Enclosure	datesale \
0	68.0	0	0	0	2006
1	4640.0	1	0	0	2004
2	2838.0	2	1	1	2004
3	3486.0	3	2	0	2011
4	722.0	2	1	2	2009
5	508.0	4	3	1	2008
6	11540.0	5	2	2	2004
7	4883.0	6	3	1	2005
8	302.0	3	2	2	2009
9	20700.0	5	0	0	2007
10	1414.0	7	3	1	2008
11	2764.0	8	4	0	2006
12	0.0	8	2	2	2005
13	1921.0	3	3	1	2006
14	320.0	1	3	1	2006
15	2450.0	9	3	1	2006
16	1972.0	5	2	2	2007
17	0.0	10	2	2	2010
18	994.0	1	2	2	2006

19	8005.0	6	3	2	2005
20	3259.0	11	2	0	2006
21	16328.0	12	5	2	2006
22	109.0	13	2	0	2007
23	17033.0	14	0	0	2006
24	1877.0	3	3	1	2004
25	1028.0	15	1	2	2006
26	356.0	16	2	2	2006
27	0.0	3	2	0	2007
28	904.0	17	5	1	2007
29	10466.0	18	2	0	2006
...
23870	9139.0	28	2	2	2011
23871	5234.0	9	2	0	2011
23872	11657.0	5	5	0	2011
23873	0.0	5	4	1	2011
23874	1032.0	41	1	1	2011
23875	11589.0	44	2	0	2011
23876	8563.0	44	3	0	2011
23877	9229.0	8	4	1	2011
23878	15060.0	8	2	2	2011
23879	17674.0	8	2	2	2011
23880	1533.0	8	2	2	2011
23881	881.0	8	3	2	2011
23882	2030.0	8	3	2	2011
23883	3333.0	1	1	1	2011
23884	5994.0	16	2	2	2011
23885	29598.0	16	0	2	2011
23886	10178.0	16	0	2	2011
23887	4607.0	16	0	2	2011
23888	0.0	16	5	2	2011
23889	23386.0	16	0	0	2011
23890	34057.0	16	0	0	2011
23891	24074.0	16	0	0	2011
23892	3138.0	16	3	0	2011
23893	1290.0	5	3	1	2011
23894	2299.0	5	0	1	2011
23895	3395.0	29	2	0	2011
23896	0.0	29	5	2	2011
23897	13507.0	29	2	2	2011
23898	5417.0	29	4	1	2011
23899	0.0	29	0	0	2011

	yearsold
0	2.0
1	8.0
2	3.0
3	10.0

4	2.0
5	4.0
6	11.0
7	4.0
8	1.0
9	1007.0
10	4.0
11	8.0
12	7.0
13	5.0
14	2.0
15	7.0
16	8.0
17	7.0
18	5.0
19	14.0
20	6.0
21	1006.0
22	2.0
23	10.0
24	3.0
25	3.0
26	3.0
27	7.0
28	3.0
29	8.0
...	...
23870	1011.0
23871	3.0
23872	11.0
23873	6.0
23874	6.0
23875	10.0
23876	6.0
23877	1011.0
23878	1011.0
23879	1011.0
23880	14.0
23881	1011.0
23882	1011.0
23883	12.0
23884	1011.0
23885	1011.0
23886	1011.0
23887	1011.0
23888	1011.0
23889	14.0
23890	1011.0

```

23891    1011.0
23892      9.0
23893      4.0
23894      7.0
23895      5.0
23896    1011.0
23897    1011.0
23898    1011.0
23899      6.0

```

```
[23900 rows x 11 columns]
```

2.1 linear

```
In [63]: regr = linear_model.LinearRegression()
```

```
In [64]: # Train the model using the training sets
regr.fit(X_train, y_train)
```

```

# Make predictions using the testing set
y_pred = regr.predict(X_test)

```

```

# The coefficients
print('Coefficients: \n', regr.coef_)
# The mean squared error
print("Mean squared error: %.2f"
      % mean_squared_error(y_test, y_pred))
# Explained variance score: 1 is perfect prediction
print('Variance score: %.2f' % r2_score(y_test, y_pred))

```

```
Coefficients:
```

```

[ 1.95134325e-03 -1.08772671e-02  1.28392054e-01  9.16600129e-13
 -1.61230572e+02  8.00475075e-01 -6.07592080e+01  3.27000956e+03
 -1.36857816e+04 -3.25624571e+02 -1.64394000e+02]

```

```
Mean squared error: 418257601.00
```

```
Variance score: 0.30
```

```
In [65]: y_pred
```

```
Out[65]: array([15308.76228324, 20361.80805096, 21858.59306422, ...,
                33147.60545152, 53148.95056587, 32171.53320709])
```

2.2 Ridge

```
In [71]: for a in [0, 0.2, 0.4, 0.6, 0.8, 1]:
```

```
    regr = Ridge(alpha=a)
```

```

# Train the model using the training sets
regr.fit(X_train, y_train)

# Make predictions using the testing set
y_pred = regr.predict(X_test)

print('alpha = ', str(a))
# The coefficients
print('Coefficients: \n', regr.coef_)
# The mean squared error
print("Mean squared error: %.2f"
      % mean_squared_error(y_test, y_pred))
# Explained variance score: 1 is perfect prediction
print('Variance score: %.2f' % r2_score(y_test, y_pred))
print('-----')

```

```

alpha = 0
Coefficients:
 [ 6.93354496e-02 -8.75989383e-03 -1.01218031e+00 -4.81631818e+16
 -3.46264392e+14 -2.03285641e+00 -6.11630211e+01  3.26988809e+03
 -1.36891390e+04  3.46264392e+14 -3.46264392e+14]
Mean squared error: 763940040.24
Variance score: -0.28
-----

```

```

alpha = 0.2
Coefficients:
 [ 1.95109156e-03 -1.08773243e-02  1.28392382e-01  0.00000000e+00
 -1.61225312e+02  8.00482577e-01 -6.07593820e+01  3.26999244e+03
 -1.36855062e+04 -3.25611526e+02 -1.64388872e+02]
Mean squared error: 418257809.51
Variance score: 0.30
-----

```

```

alpha = 0.4
Coefficients:
 [ 1.95083988e-03 -1.08773815e-02  1.28392709e-01  0.00000000e+00
 -1.61218254e+02  8.00490078e-01 -6.07595560e+01  3.26997531e+03
 -1.36852309e+04 -3.25600280e+02 -1.64381944e+02]
Mean squared error: 418258018.10
Variance score: 0.30
-----

```

```

alpha = 0.6
Coefficients:
 [ 1.95058822e-03 -1.08774387e-02  1.28393036e-01  0.00000000e+00
 -1.61212335e+02  8.00497578e-01 -6.07597300e+01  3.26995819e+03
 -1.36849556e+04 -3.25587895e+02 -1.64376157e+02]
Mean squared error: 418258226.76
Variance score: 0.30
-----

```

```
alpha = 0.8
Coefficients:
[ 1.95033657e-03 -1.08774959e-02  1.28393364e-01  0.00000000e+00
 -1.61206109e+02  8.00505079e-01 -6.07599039e+01  3.26994106e+03
 -1.36846803e+04 -3.25575819e+02 -1.64370063e+02]
Mean squared error: 418258435.50
Variance score: 0.30
-----
```

```
alpha = 1
Coefficients:
[ 1.95008493e-03 -1.08775531e-02  1.28393691e-01  0.00000000e+00
 -1.61199937e+02  8.00512579e-01 -6.07600779e+01  3.26992394e+03
 -1.36844050e+04 -3.25563689e+02 -1.64364022e+02]
Mean squared error: 418258644.31
Variance score: 0.30
-----
```

```
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/ridge.py:
Ill-conditioned matrix detected. Result is not guaranteed to be accurate.
Reciprocal condition number1.061562e-16
  overwrite_a=True).T
```

```
In [90]: for a in [0.18, 0.19, 0.2, 0.21, 0.22, 0.23]:
```

```
    regr = Ridge(alpha=a)

    # Train the model using the training sets
    regr.fit(X_train, y_train)

    # Make predictions using the testing set
    y_pred = regr.predict(X_test)

    print('alpha = ', str(a))
    # The coefficients
    print('Coefficients: \n', regr.coef_)
    # The mean squared error
    print("Mean squared error: %.2f"
          % mean_squared_error(y_test, y_pred))
    # Explained variance score: 1 is perfect prediction
    print('Variance score: %.2f' % r2_score(y_test, y_pred))
    print('-----')
```

```
alpha = 0.18
Coefficients:
[ 1.95111673e-03 -1.08773186e-02  1.28392349e-01  0.00000000e+00
 -1.61225920e+02  8.00481827e-01 -6.07593646e+01  3.26999415e+03
```



```
-1.36855337e+04 -3.25612749e+02 -1.64389466e+02]
```

```
Mean squared error: 418257788.66
```

```
Variance score: 0.30
```

```
-----
```

```
alpha = 0.19
```

```
Coefficients:
```

```
[ 1.95110414e-03 -1.08773214e-02  1.28392365e-01  0.00000000e+00
```

```
-1.61225339e+02  8.00482202e-01 -6.07593733e+01  3.26999329e+03
```

```
-1.36855200e+04 -3.25612414e+02 -1.64388892e+02]
```

```
Mean squared error: 418257799.08
```

```
Variance score: 0.30
```

```
-----
```

```
alpha = 0.2
```

```
Coefficients:
```

```
[ 1.95109156e-03 -1.08773243e-02  1.28392382e-01  0.00000000e+00
```

```
-1.61225312e+02  8.00482577e-01 -6.07593820e+01  3.26999244e+03
```

```
-1.36855062e+04 -3.25611526e+02 -1.64388872e+02]
```

```
Mean squared error: 418257809.51
```

```
Variance score: 0.30
```

```
-----
```

```
alpha = 0.21
```

```
Coefficients:
```

```
[ 1.95107898e-03 -1.08773271e-02  1.28392398e-01  0.00000000e+00
```

```
-1.61224751e+02  8.00482952e-01 -6.07593907e+01  3.26999158e+03
```

```
-1.36854924e+04 -3.25611172e+02 -1.64388317e+02]
```

```
Mean squared error: 418257819.94
```

```
Variance score: 0.30
```

```
-----
```

```
alpha = 0.22
```

```
Coefficients:
```

```
[ 1.95106639e-03 -1.08773300e-02  1.28392414e-01  0.00000000e+00
```

```
-1.61224559e+02  8.00483327e-01 -6.07593994e+01  3.26999072e+03
```

```
-1.36854787e+04 -3.25610448e+02 -1.64388132e+02]
```

```
Mean squared error: 418257830.37
```

```
Variance score: 0.30
```

```
-----
```

```
alpha = 0.23
```

```
Coefficients:
```

```
[ 1.95105381e-03 -1.08773329e-02  1.28392431e-01  0.00000000e+00
```

```
-1.61223914e+02  8.00483702e-01 -6.07594081e+01  3.26998987e+03
```

```
-1.36854649e+04 -3.25610179e+02 -1.64387493e+02]
```

```
Mean squared error: 418257840.79
```

```
Variance score: 0.30
```

```
-----
```

```
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/ridge.py:
```

```
Ill-conditioned matrix detected. Result is not guaranteed to be accurate.
```

```

Reciprocal condition number9.554057e-17
  overwrite_a=True).T
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/ridge.py:
Ill-conditioned matrix detected. Result is not guaranteed to be accurate.
Reciprocal condition number1.008485e-16
  overwrite_a=True).T
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/ridge.py:
Ill-conditioned matrix detected. Result is not guaranteed to be accurate.
Reciprocal condition number1.061562e-16
  overwrite_a=True).T

```

```

In [92]: for a in [0.1, 0.11, 0.12, 0.13, 0.14, 0.15, 0.16, 0.7]:

```

```

    regr = Ridge(alpha=a)

    # Train the model using the training sets
    regr.fit(X_train, y_train)

    # Make predictions using the testing set
    y_pred = regr.predict(X_test)

    print('alpha = ', str(a))
    # The coefficients
    print('Coefficients: \n', regr.coef_)
    # The mean squared error
    print("Mean squared error: %.2f"
          % mean_squared_error(y_test, y_pred))
    # Explained variance score: 1 is perfect prediction
    print('Variance score: %.2f' % r2_score(y_test, y_pred))
    print('-----')

```

```

alpha = 0.1
Coefficients:
 [ 1.95121740e-03 -1.08772957e-02  1.28392218e-01  0.00000000e+00
 -1.61228856e+02  8.00478826e-01 -6.07592950e+01  3.27000100e+03
 -1.36856439e+04 -3.25617135e+02 -1.64392349e+02]
Mean squared error: 418257705.25
Variance score: 0.30
-----

```

```

alpha = 0.11
Coefficients:
 [ 1.95120482e-03 -1.08772985e-02  1.28392234e-01  0.00000000e+00
 -1.61228053e+02  8.00479201e-01 -6.07593037e+01  3.27000014e+03
 -1.36856301e+04 -3.25617023e+02 -1.64391553e+02]
Mean squared error: 418257715.67
Variance score: 0.30
-----

```

```

alpha = 0.12
Coefficients:
 [ 1.95119223e-03 -1.08773014e-02  1.28392251e-01  0.00000000e+00
 -1.61227695e+02  8.00479576e-01 -6.07593124e+01  3.26999929e+03
 -1.36856163e+04 -3.25616465e+02 -1.64391202e+02]
Mean squared error: 418257726.10
Variance score: 0.30
-----

alpha = 0.13
Coefficients:
 [ 1.95117965e-03 -1.08773043e-02  1.28392267e-01  0.00000000e+00
 -1.61227165e+02  8.00479951e-01 -6.07593211e+01  3.26999843e+03
 -1.36856026e+04 -3.25616079e+02 -1.64390679e+02]
Mean squared error: 418257736.52
Variance score: 0.30
-----

alpha = 0.14
Coefficients:
 [ 1.95116707e-03 -1.08773071e-02  1.28392283e-01  0.00000000e+00
 -1.61227241e+02  8.00480326e-01 -6.07593298e+01  3.26999757e+03
 -1.36855888e+04 -3.25615089e+02 -1.64390761e+02]
Mean squared error: 418257746.95
Variance score: 0.30
-----

alpha = 0.15
Coefficients:
 [ 1.95115448e-03 -1.08773100e-02  1.28392300e-01  0.00000000e+00
 -1.61226739e+02  8.00480701e-01 -6.07593385e+01  3.26999672e+03
 -1.36855750e+04 -3.25614675e+02 -1.64390265e+02]
Mean squared error: 418257757.38
Variance score: 0.30
-----

alpha = 0.16
Coefficients:
 [ 1.95114190e-03 -1.08773128e-02  1.28392316e-01  0.00000000e+00
 -1.61226448e+02  8.00481076e-01 -6.07593472e+01  3.26999586e+03
 -1.36855613e+04 -3.25614051e+02 -1.64389981e+02]
Mean squared error: 418257767.80
Variance score: 0.30
-----

alpha = 0.7
Coefficients:
 [ 1.95046240e-03 -1.08774673e-02  1.28393200e-01  0.00000000e+00
 -1.61209216e+02  8.00501329e-01 -6.07598169e+01  3.26994963e+03
 -1.36848179e+04 -3.25581864e+02 -1.64373104e+02]
Mean squared error: 418258331.12
Variance score: 0.30
-----

```

```

/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/ridge.py:
Ill-conditioned matrix detected. Result is not guaranteed to be accurate.
Reciprocal condition number5.307793e-17
    overwrite_a=True).T
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/ridge.py:
Ill-conditioned matrix detected. Result is not guaranteed to be accurate.
Reciprocal condition number5.838591e-17
    overwrite_a=True).T
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/ridge.py:
Ill-conditioned matrix detected. Result is not guaranteed to be accurate.
Reciprocal condition number6.369374e-17
    overwrite_a=True).T
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/ridge.py:
Ill-conditioned matrix detected. Result is not guaranteed to be accurate.
Reciprocal condition number6.900165e-17
    overwrite_a=True).T
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/ridge.py:
Ill-conditioned matrix detected. Result is not guaranteed to be accurate.
Reciprocal condition number7.430929e-17
    overwrite_a=True).T
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/ridge.py:
Ill-conditioned matrix detected. Result is not guaranteed to be accurate.
Reciprocal condition number7.961719e-17
    overwrite_a=True).T
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/ridge.py:
Ill-conditioned matrix detected. Result is not guaranteed to be accurate.
Reciprocal condition number8.492500e-17
    overwrite_a=True).T

```

2.3 Lasso

```
In [72]: for a in [0, 0.2, 0.4, 0.6, 0.8, 1]:
```

```

    regr = linear_model.Lasso(alpha=a)

    # Train the model using the training sets
    regr.fit(X_train, y_train)

    # Make predictions using the testing set
    y_pred = regr.predict(X_test)

    print('alpha = ', str(a))
    # The coefficients
    print('Coefficients: \n', regr.coef_)
    # The mean squared error

```

```

print("Mean squared error: %.2f"
      % mean_squared_error(y_test, y_pred))
# Explained variance score: 1 is perfect prediction
print('Variance score: %.2f' % r2_score(y_test, y_pred))
print('-----')

```

```

/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/ipykernel_launcher.py:6: UserWarning:
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/coordinat
positive)
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/coordinat
ConvergenceWarning)

```

```

alpha = 0
Coefficients:
[ 1.95134325e-03 -1.08772671e-02  1.28392054e-01  0.00000000e+00
  5.27928021e+00  8.00475075e-01 -6.07592080e+01  3.27000956e+03
 -1.36857816e+04 -4.92134423e+02  2.11585219e+00]
Mean squared error: 418257601.00
Variance score: 0.30
-----

```

```

alpha = 0.2
Coefficients:
[ 1.94732410e-03 -1.08774212e-02  1.28389354e-01  0.00000000e+00
  5.27514744e+00  8.00482359e-01 -6.07582856e+01  3.26991438e+03
 -1.36853732e+04 -4.92030133e+02  2.11148863e+00]
Mean squared error: 418257744.51
Variance score: 0.30
-----

```

```

alpha = 0.4
Coefficients:
[ 1.94330495e-03 -1.08775753e-02  1.28386655e-01  0.00000000e+00
  5.27101468e+00  8.00489643e-01 -6.07573633e+01  3.26981920e+03
 -1.36849648e+04 -4.91925842e+02  2.10712506e+00]
Mean squared error: 418257888.28
Variance score: 0.30
-----

```

```

alpha = 0.6
Coefficients:
[ 1.93928581e-03 -1.08777293e-02  1.28383955e-01  0.00000000e+00
  5.26688192e+00  8.00496926e-01 -6.07564409e+01  3.26972403e+03
 -1.36845564e+04 -4.91821552e+02  2.10276150e+00]
Mean squared error: 418258032.28
Variance score: 0.30
-----

```

```

alpha = 0.8
Coefficients:

```

```

[ 1.93526666e-03 -1.08778834e-02  1.28381255e-01  0.00000000e+00
 5.26274916e+00  8.00504210e-01 -6.0755186e+01  3.26962885e+03
-1.36841480e+04 -4.91717262e+02  2.09839794e+00]
Mean squared error: 418258176.54
Variance score: 0.30
-----
alpha = 1
Coefficients:
[ 1.93124751e-03 -1.08780375e-02  1.28378555e-01  0.00000000e+00
 5.25861640e+00  8.00511494e-01 -6.07545962e+01  3.26953367e+03
-1.36837396e+04 -4.91612971e+02  2.09403437e+00]
Mean squared error: 418258321.05
Variance score: 0.30
-----

```

2.4 Elastic Net

```
In [73]: for a in [0, 0.2, 0.4, 0.6, 0.8, 1]:
```

```

    regr = linear_model.ElasticNet(alpha=a)

    # Train the model using the training sets
    regr.fit(X_train, y_train)

    # Make predictions using the testing set
    y_pred = regr.predict(X_test)

    print('alpha = ', str(a))
    # The coefficients
    print('Coefficients: \n', regr.coef_)
    # The mean squared error
    print("Mean squared error: %.2f"
          % mean_squared_error(y_test, y_pred))
    # Explained variance score: 1 is perfect prediction
    print('Variance score: %.2f' % r2_score(y_test, y_pred))
    print('-----')

```

```
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/ipykernel_launcher.py:6: UserWarning
```

```
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:100: ConvergenceWarning
```

```
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:100: ConvergenceWarning
```

```

alpha = 0
Coefficients:

```

```

[ 1.95134325e-03 -1.08772671e-02  1.28392054e-01  0.00000000e+00
 5.27928021e+00  8.00475075e-01 -6.07592080e+01  3.27000956e+03
-1.36857816e+04 -4.92134423e+02  2.11585219e+00]
Mean squared error: 418257601.00
Variance score: 0.30
-----
alpha = 0.2
Coefficients:
[ -1.81447077e-04 -1.13441265e-02  1.30691043e-01  0.00000000e+00
  4.96660340e+00  8.60001910e-01 -6.21619769e+01  3.11522203e+03
 -1.14770290e+04 -3.42362096e+02  7.54373158e-01]
Mean squared error: 422599616.41
Variance score: 0.29
-----
alpha = 0.4
Coefficients:
[ -1.47706023e-03 -1.16885398e-02  1.32082847e-01  0.00000000e+00
  4.82008485e+00  9.02091622e-01 -6.31952279e+01  2.97598486e+03
 -9.88272856e+03 -2.38534440e+02 -1.35216045e-01]
Mean squared error: 429090896.23
Variance score: 0.28
-----
alpha = 0.6
Coefficients:
[ -2.26474672e-03 -1.19545218e-02  1.32930982e-01  0.00000000e+00
  4.77287391e+00  9.33229136e-01 -6.39909455e+01  2.84957205e+03
 -8.67774452e+03 -1.63433171e+02 -7.33239085e-01]
Mean squared error: 435900056.01
Variance score: 0.27
-----
alpha = 0.8
Coefficients:
[ -2.72767242e-03 -1.21670491e-02  1.33436124e-01  0.00000000e+00
  4.78413259e+00  9.57062363e-01 -6.46241212e+01  2.73402719e+03
 -7.73496058e+03 -1.07378280e+02 -1.14440228e+00]
Mean squared error: 442398051.36
Variance score: 0.26
-----
alpha = 1
Coefficients:
[ -2.97504431e-03 -1.23413818e-02  1.33716788e-01  0.00000000e+00
  4.83343402e+00  9.75796959e-01 -6.51406469e+01  2.62785703e+03
 -6.97717129e+03 -6.45283502e+01 -1.42778011e+00]
Mean squared error: 448388144.43
Variance score: 0.25
-----

```

```
In [74]: for a in [0, 0.2, 0.4, 0.6, 0.8, 1]:
```

```
    regr = linear_model.ElasticNet(alpha=0, l1_ratio=a)

    # Train the model using the training sets
    regr.fit(X_train, y_train)

    # Make predictions using the testing set
    y_pred = regr.predict(X_test)

    print('alpha = ', str(a))
    # The coefficients
    print('Coefficients: \n', regr.coef_)
    # The mean squared error
    print("Mean squared error: %.2f"
          % mean_squared_error(y_test, y_pred))
    # Explained variance score: 1 is perfect prediction
    print('Variance score: %.2f' % r2_score(y_test, y_pred))
    print('-----')
```

/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/ipykernel_launcher.py:6: UserWarning

/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/coordinate_descent.py:111: ConvergenceWarning

```
alpha = 0
Coefficients:
[ 1.95134325e-03 -1.08772671e-02  1.28392054e-01  0.00000000e+00
  5.27928021e+00  8.00475075e-01 -6.07592080e+01  3.27000956e+03
 -1.36857816e+04 -4.92134423e+02  2.11585219e+00]
Mean squared error: 418257601.00
Variance score: 0.30
-----
```

```
alpha = 0.2
Coefficients:
[ 1.95134325e-03 -1.08772671e-02  1.28392054e-01  0.00000000e+00
  5.27928021e+00  8.00475075e-01 -6.07592080e+01  3.27000956e+03
 -1.36857816e+04 -4.92134423e+02  2.11585219e+00]
Mean squared error: 418257601.00
Variance score: 0.30
-----
```

```
alpha = 0.4
Coefficients:
[ 1.95134325e-03 -1.08772671e-02  1.28392054e-01  0.00000000e+00
  5.27928021e+00  8.00475075e-01 -6.07592080e+01  3.27000956e+03
```



```

-1.36857816e+04 -4.92134423e+02  2.11585219e+00]
Mean squared error: 418257601.00
Variance score: 0.30
-----
alpha = 0.6
Coefficients:
[ 1.95134325e-03 -1.08772671e-02  1.28392054e-01  0.00000000e+00
 5.27928021e+00  8.00475075e-01 -6.07592080e+01  3.27000956e+03
-1.36857816e+04 -4.92134423e+02  2.11585219e+00]
Mean squared error: 418257601.00
Variance score: 0.30
-----
alpha = 0.8
Coefficients:
[ 1.95134325e-03 -1.08772671e-02  1.28392054e-01  0.00000000e+00
 5.27928021e+00  8.00475075e-01 -6.07592080e+01  3.27000956e+03
-1.36857816e+04 -4.92134423e+02  2.11585219e+00]
Mean squared error: 418257601.00
Variance score: 0.30
-----
alpha = 1
Coefficients:
[ 1.95134325e-03 -1.08772671e-02  1.28392054e-01  0.00000000e+00
 5.27928021e+00  8.00475075e-01 -6.07592080e+01  3.27000956e+03
-1.36857816e+04 -4.92134423e+02  2.11585219e+00]
Mean squared error: 418257601.00
Variance score: 0.30
-----

```

```

In [76]: for a in [3, 10, 18, 50, 100, 1000]:

    regr = linear_model.ElasticNet(alpha=0, max_iter=a)

    # Train the model using the training sets
    regr.fit(X_train, y_train)

    # Make predictions using the testing set
    y_pred = regr.predict(X_test)

    print('alpha = ', str(a))
    # The coefficients
    print('Coefficients: \n', regr.coef_)
    # The mean squared error
    print("Mean squared error: %.2f"
          % mean_squared_error(y_test, y_pred))
    # Explained variance score: 1 is perfect prediction
    print('Variance score: %.2f' % r2_score(y_test, y_pred))

```

```

print('-----')

/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/ipykernel_launcher.py:6: UserWarning:
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/coordinat
positive)
/Users/chrispaul/anaconda2/envs/nlp/lib/python3.6/site-packages/sklearn/linear_model/coordinat
ConvergenceWarning)

alpha = 3
Coefficients:
[-9.31445297e-04 -1.09747225e-02  1.31362963e-01  0.00000000e+00
 5.23783572e+00  8.07259837e-01 -6.06477363e+01  3.27237151e+03
-1.36469352e+04 -4.65112075e+02  2.04766015e+00]
Mean squared error: 418427184.07
Variance score: 0.30
-----
alpha = 10
Coefficients:
[ 1.94288381e-03 -1.08773894e-02  1.28388620e-01  0.00000000e+00
 5.27922274e+00  8.00477128e-01 -6.07590793e+01  3.27000972e+03
-1.36857166e+04 -4.92061726e+02  2.11573827e+00]
Mean squared error: 418257650.54
Variance score: 0.30
-----
alpha = 18
Coefficients:
[ 1.95133561e-03 -1.08772672e-02  1.28392051e-01  0.00000000e+00
 5.27928015e+00  8.00475077e-01 -6.07592078e+01  3.27000956e+03
-1.36857815e+04 -4.92134357e+02  2.11585209e+00]
Mean squared error: 418257601.04
Variance score: 0.30
-----
alpha = 50
Coefficients:
[ 1.95134325e-03 -1.08772671e-02  1.28392054e-01  0.00000000e+00
 5.27928021e+00  8.00475075e-01 -6.07592080e+01  3.27000956e+03
-1.36857816e+04 -4.92134423e+02  2.11585219e+00]
Mean squared error: 418257601.00
Variance score: 0.30
-----
alpha = 100
Coefficients:
[ 1.95134325e-03 -1.08772671e-02  1.28392054e-01  0.00000000e+00
 5.27928021e+00  8.00475075e-01 -6.07592080e+01  3.27000956e+03
-1.36857816e+04 -4.92134423e+02  2.11585219e+00]
Mean squared error: 418257601.00

```

Variance score: 0.30

alpha = 1000

Coefficients:

```
[ 1.95134325e-03 -1.08772671e-02  1.28392054e-01  0.00000000e+00
  5.27928021e+00  8.00475075e-01 -6.07592080e+01  3.27000956e+03
 -1.36857816e+04 -4.92134423e+02  2.11585219e+00]
```

Mean squared error: 418257601.00

Variance score: 0.30

2.5 RF

```
In [81]: parameters = []
        mse = []
```

```
    for n in [2, 5, 10, 20, 50, 100]:
        for md in [3, 5, 10, 20, 50, 100, 1000]:
            for mss in [2, 3, 5, 8, 10, 20]:
```

```
                regr = RandomForestRegressor(max_depth=md, n_estimators=n, min_samples_sp
```

```
                regr.fit(X_train, y_train)
```

```
                y_pred = regr.predict(X_test)
```

```
                mse_ = np.sqrt(np.mean((y_pred - y_test)**2))
```

```
                parameters.append("n_estimators = " + str(n)
                                   + ", max_depth = " + str(md)
                                   + ", min_samples_split = " + str(mss))
```

```
                mse.append(mse_)
```

```
    print("a sixth")
```

a sixth

a sixth

a sixth

a sixth

a sixth

a sixth

```
In [85]: ind = np.argmin(mse)
```

```
In [86]: ind
```

```
Out [86]: 236
```

```
In [87]: mse[236]
```

```
Out[87]: 9758.55780823173
```

```
In [89]: parameters[236]
```

```
Out[89]: 'n_estimators = 100, max_depth = 50, min_samples_split = 5'
```